

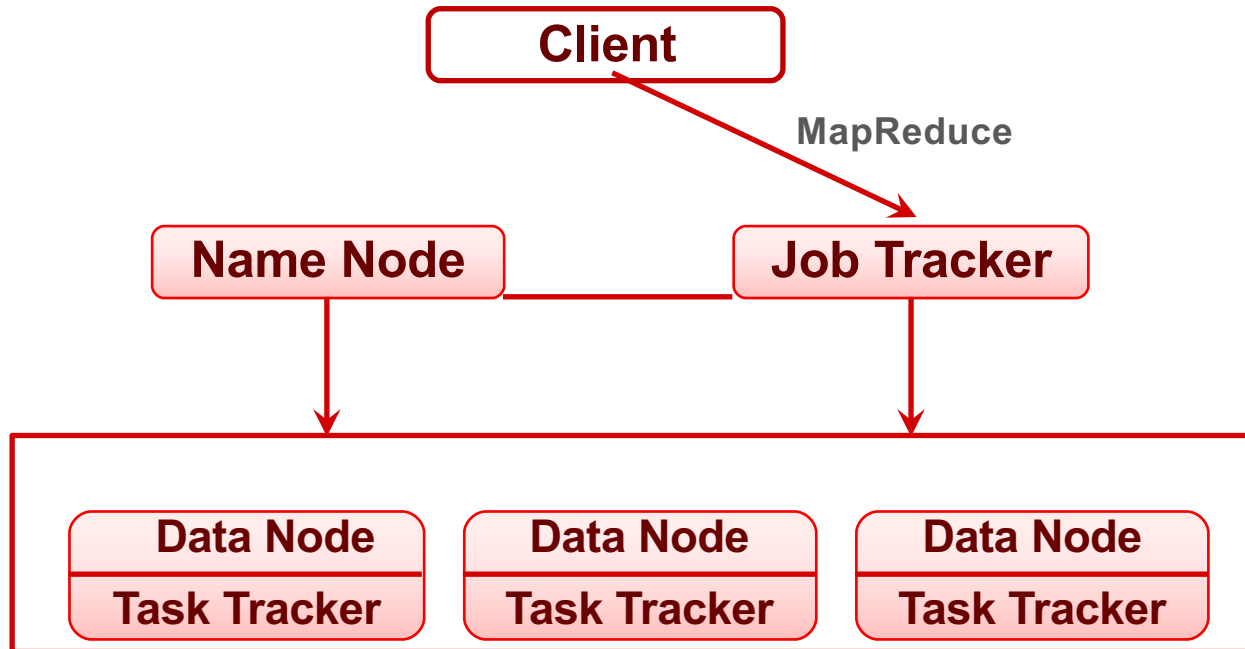



Module 4_2

Understanding MapReduce and Ozzie

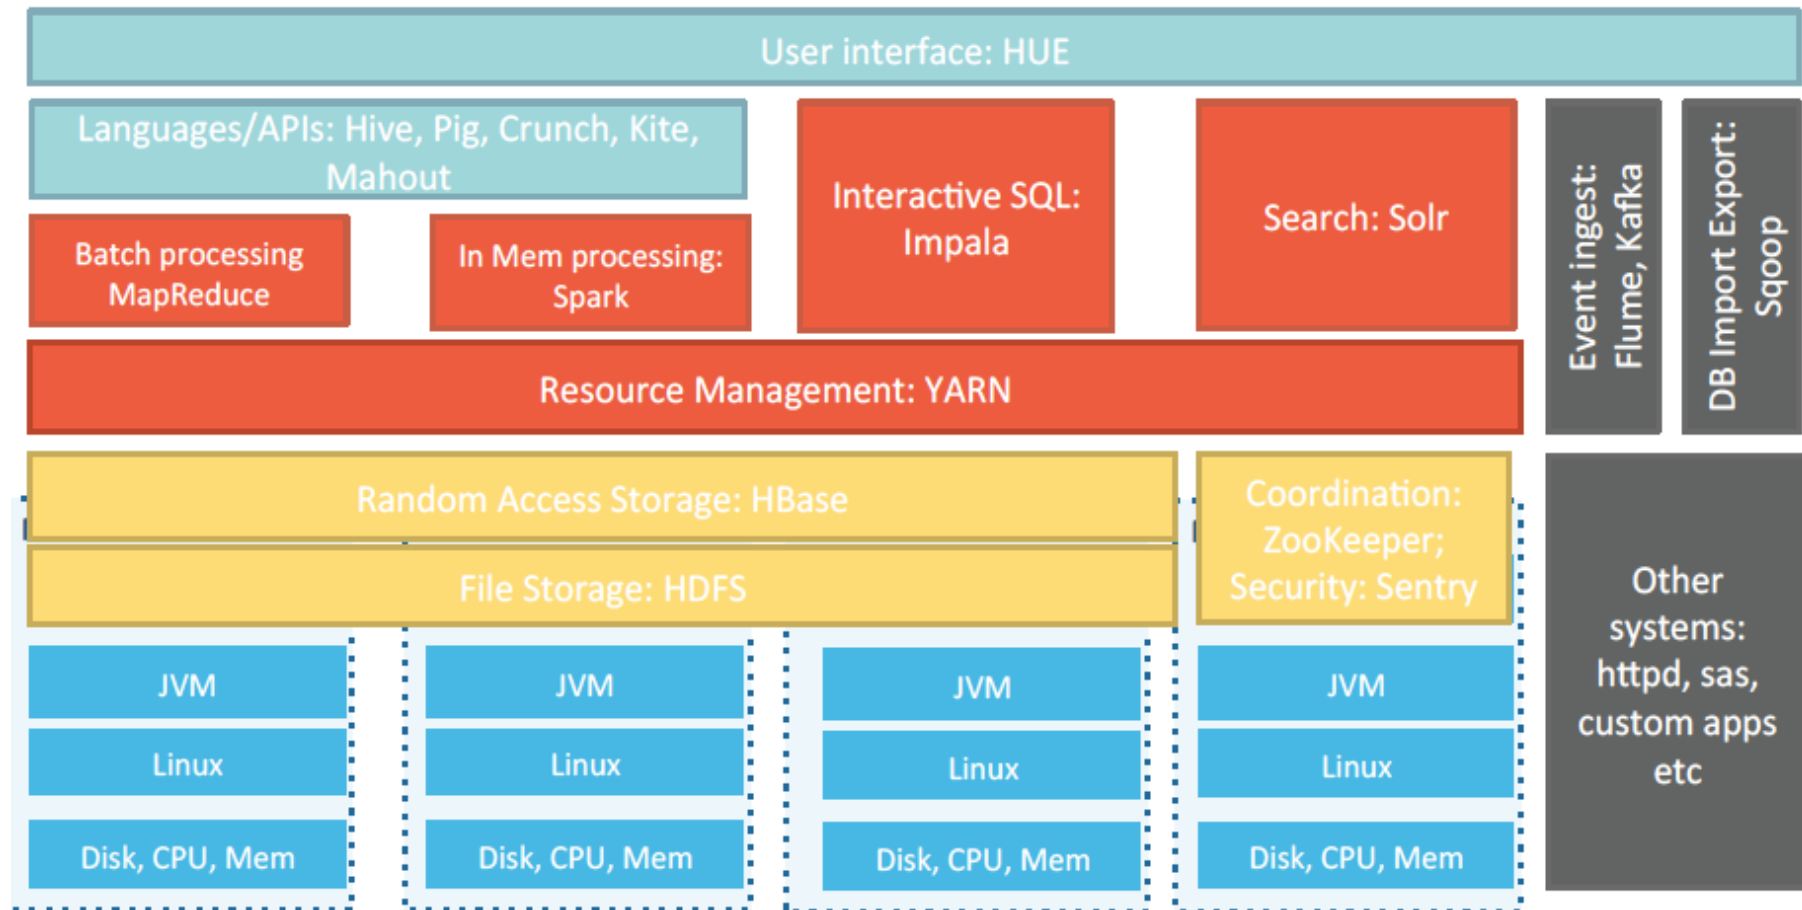
Thanachart Numnonda, Executive Director, IMC Institute

Thanisa Numnonda, Faculty of Information Technology,
King Mongkut's Institute of Technology Ladkrabang

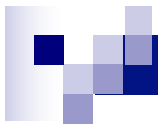


Lecture: Understanding MapReduce Processing

Hadoop Ecosystem



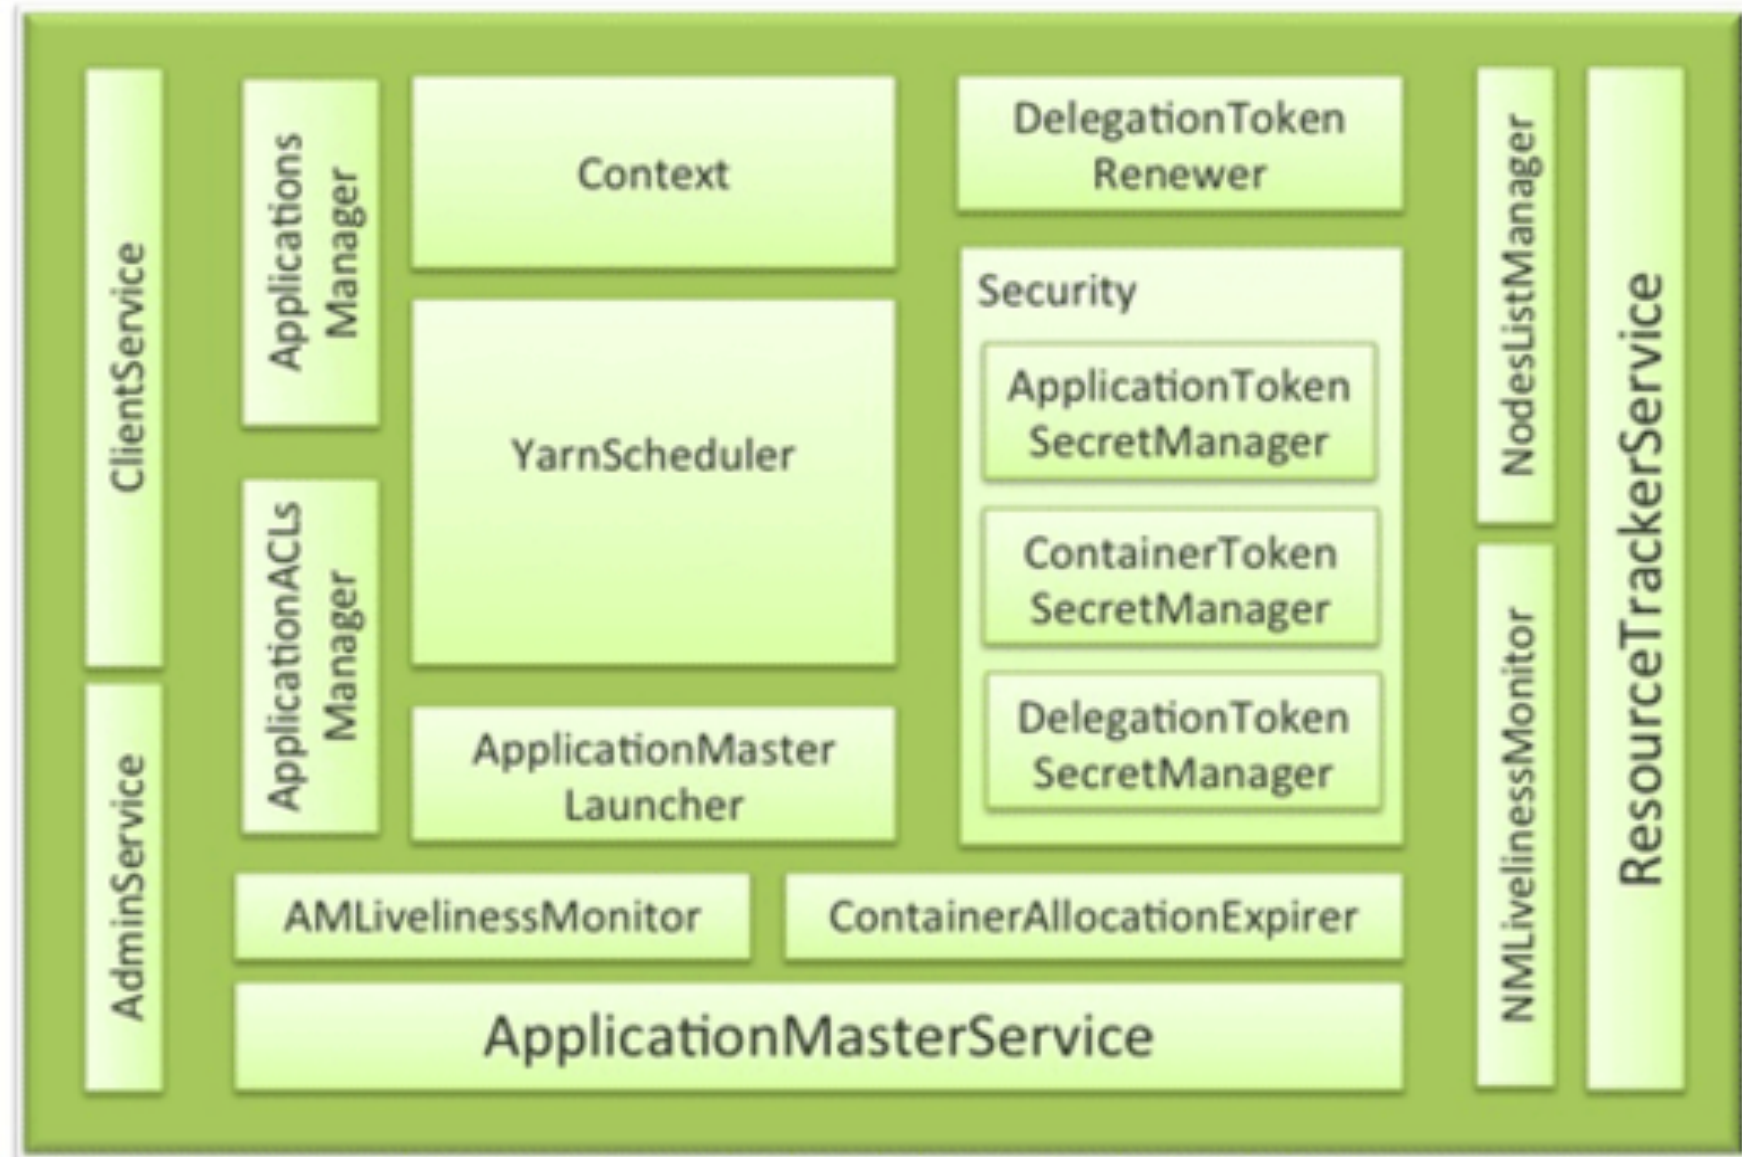
Source: Apache Hadoop Operations for Production Systems, Cloudera



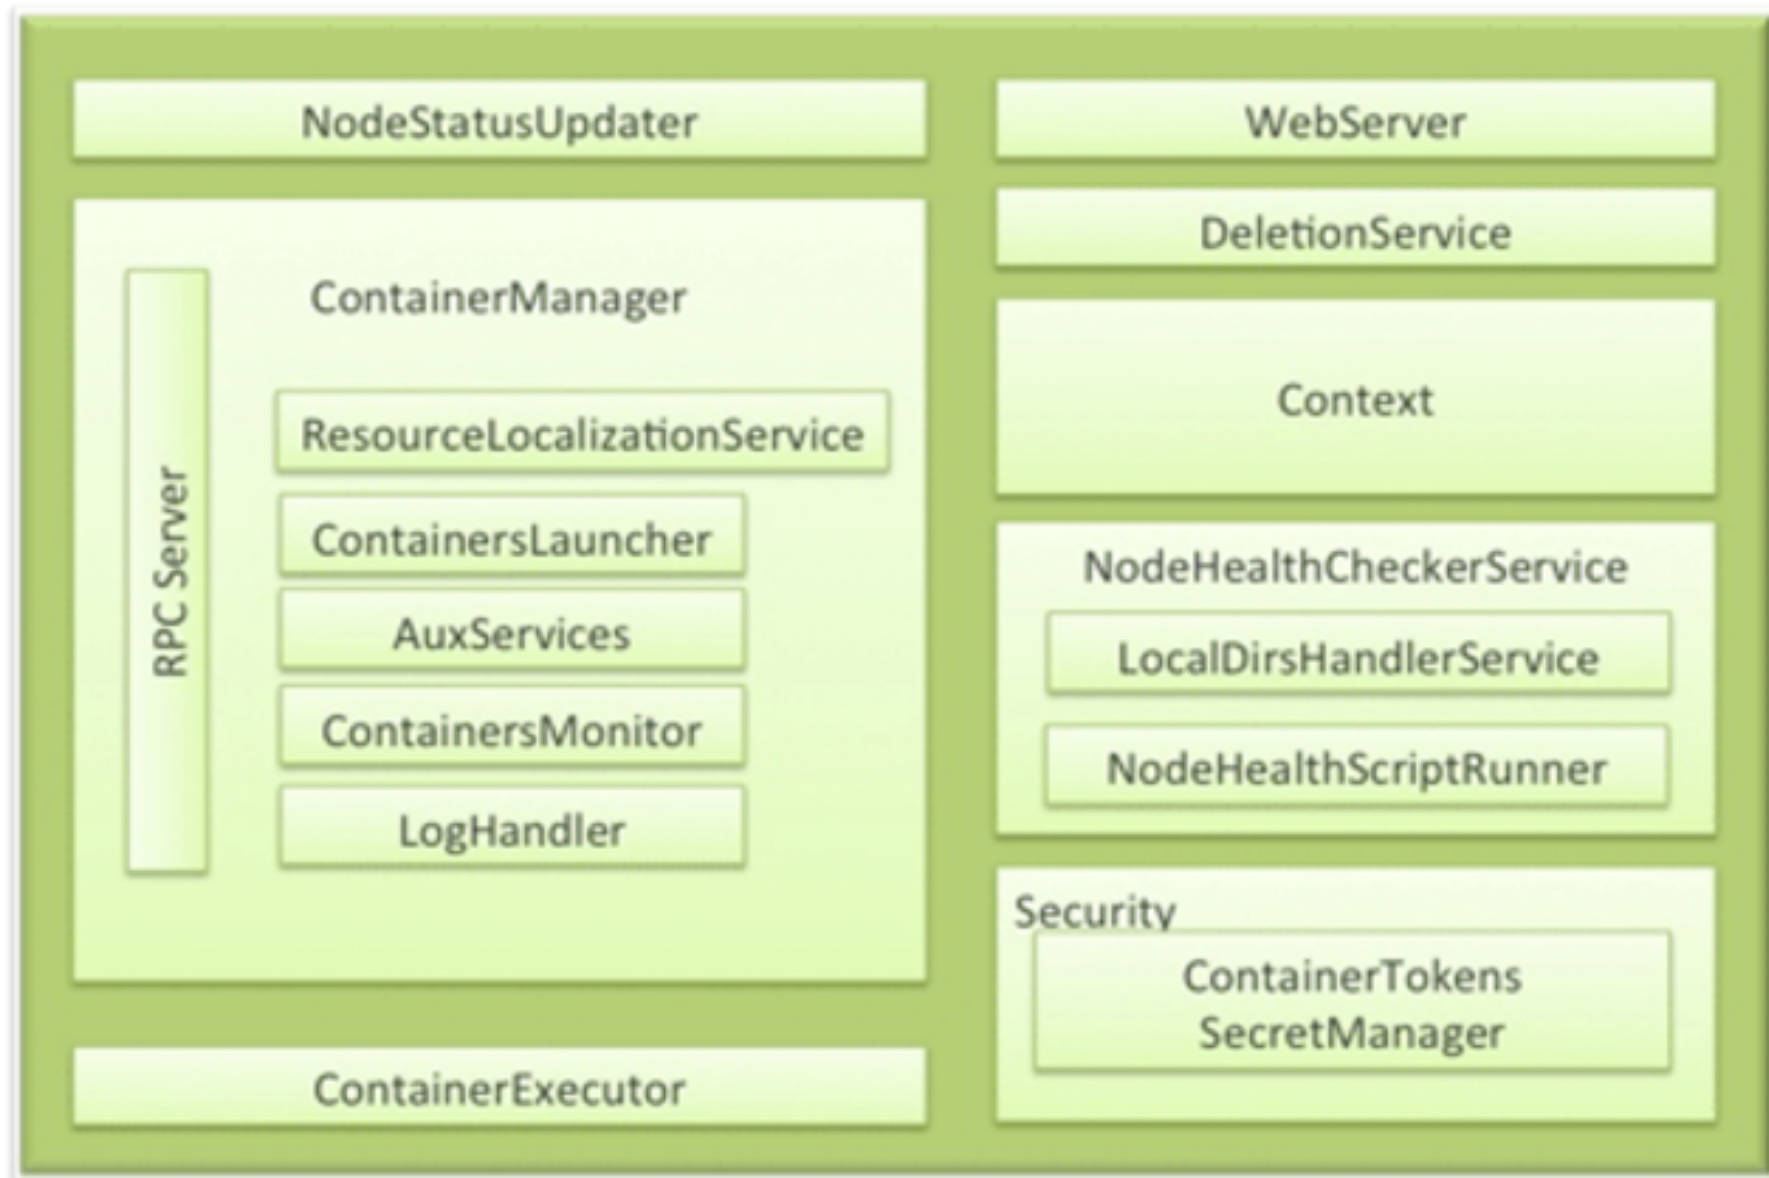
Yarn Components

- Resource Manager
 - a pluggable scheduler > primarily limited to scheduling
 - ApplicationManager >> manages user jobs on the cluster
- Node Manager
 - Containers
 - manages users' jobs and workflow on a given node
- Application Master
 - a user job life-cycle manager


Resource Manager



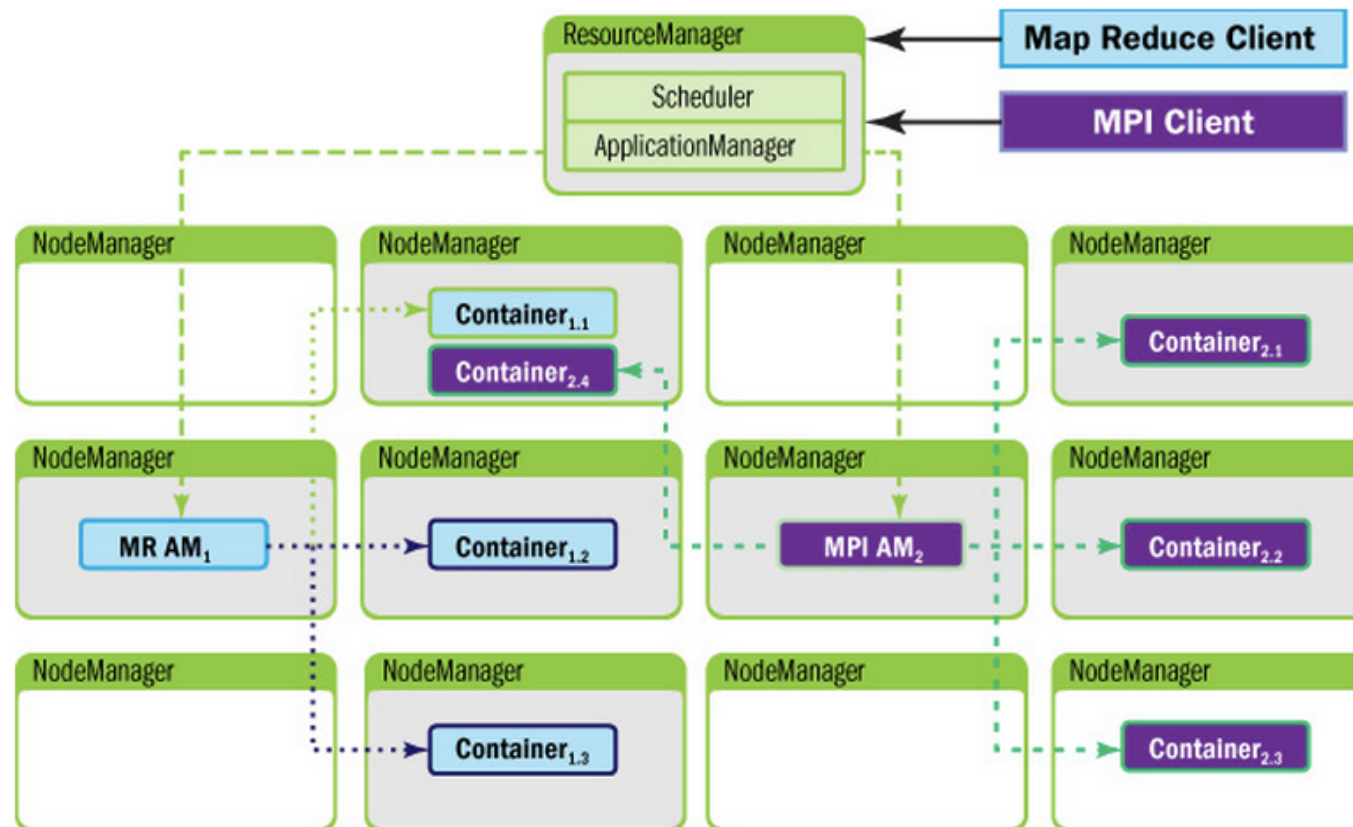
Node Manager



YARN: – CONCEPTS APPLICATIONS



YARN: – CONCEPTS APPLICATIONS





Yarn Workloads

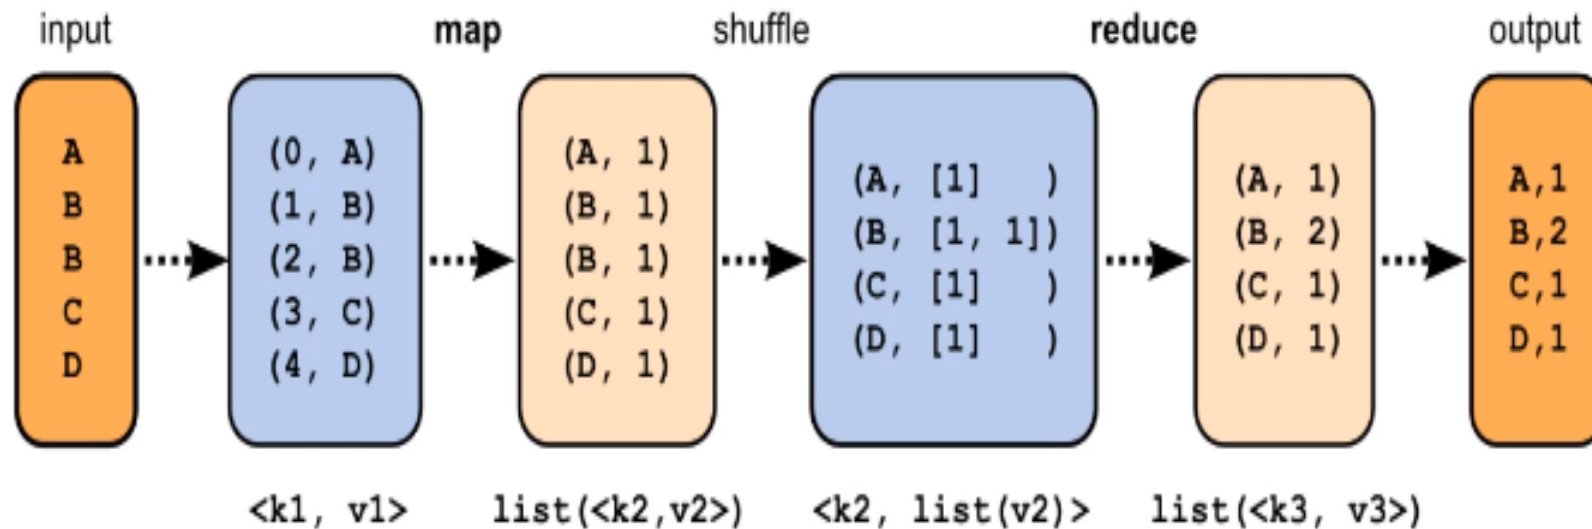
- **Batch:** MapReduce that is the compatible with Hadoop 1.x
- **Script:** Pig
- **Interactive SQL:** Hive or Tez
- **NoSQL:** HBase and Accumulo
- **Streaming:** Storm
- **In-memory:** Spark
- **Search:** SOLR

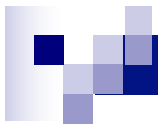


Before MapReduce...

- Large scale data processing was difficult!
 - Managing hundreds or thousands of processors
 - Managing parallelization and distribution
 - I/O Scheduling
 - Status and monitoring
 - Fault/crash tolerance
- MapReduce provides all of these, easily!

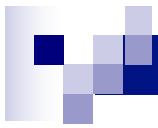
MapReduce Framework






How Map and Reduce Work Together

- Map returns information
- Reduces accepts information
- Reduce applies a user defined function to reduce the amount of data



Map Abstraction

- Inputs a key/value pair
 - Key is a reference to the input value
 - Value is the data set on which to operate
- Evaluation
 - Function defined by user
 - Applies to every value in value input
 - Might need to parse input
- Produces a new list of key/value pairs
 - Can be different type from input pair



Reduce Abstraction

- Starts with intermediate Key / Value pairs
- Ends with finalized Key / Value pairs
- Starting pairs are sorted by key
- Iterator supplies the values for a given key to the Reduce function.



Reduce Abstraction

- Typically a function that:
 - Starts with a large number of key/value pairs
 - One key/value for each word in all files being greped (including multiple entries for the same word)
 - Ends with very few key/value pairs
 - One key/value for each unique word across all the files with the number of instances summed into this entry
- Broken up so a given worker works with input of the same key.



Why is this approach better?

- Creates an abstraction for dealing with complex overhead
 - The computations are simple, the overhead is messy
- Removing the overhead makes programs much smaller and thus easier to use
 - Less testing is required as well. The MapReduce libraries can be assumed to work properly, so only user code needs to be tested
- Division of labor also handled by the MapReduce libraries, so programmers only need to focus on the actual computation



Hands-On: Writing you MapReduce Program



Example MapReduce: WordCount

```
package org.apache.hadoop.examples;

import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.util.GenericOptionsParser;

public class WordCount {

    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable>{

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context
            ) throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }
}
```

```

public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values,
        Context context
        ) throws IOException, InterruptedException {

        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}

```

```

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    String[] otherArgs = new GenericOptionsParser(conf, args).getRemainingArgs();
    if (otherArgs.length != 2) {
        System.err.println("Usage: wordcount <in> <out>");
        System.exit(2);
    }
    Job job = new Job(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(otherArgs[0]));
    FileOutputFormat.setOutputPath(job, new Path(otherArgs[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```



Running MapReduce Program

```
[cloudera@quickstart ~]$ cd workspace/
[cloudera@quickstart workspace]$ hadoop jar wordcount.jar org.myorg.WordCount input/* output/wordcount_output
15/02/08 10:30:31 INFO client.RMPProxy: Connecting to ResourceManager at /0.0.0.0:8032
15/02/08 10:30:32 INFO client.RMPProxy: Connecting to ResourceManager at /0.0.0.0:8032
15/02/08 10:30:33 WARN mapreduce.JobSubmitter: Hadoop command-line option parsing not performed. Implement the Tool i
ation with ToolRunner to remedy this.
15/02/08 10:30:33 INFO mapred.FileInputFormat: Total input paths to process : 1
15/02/08 10:30:34 INFO mapreduce.JobSubmitter: number of splits:2
15/02/08 10:30:34 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1423408479621_0009
15/02/08 10:30:35 INFO impl.YarnClientImpl: Submitted application application_1423408479621_0009
15/02/08 10:30:35 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_142
15/02/08 10:30:35 INFO mapreduce.Job: Running job: job_1423408479621_0009
15/02/08 10:30:52 INFO mapreduce.Job: Job job_1423408479621_0009 running in uber mode : false
15/02/08 10:30:52 INFO mapreduce.Job: map 0% reduce 0%
15/02/08 10:31:22 INFO mapreduce.Job: map 58% reduce 0%
15/02/08 10:31:25 INFO mapreduce.Job: map 100% reduce 0%
15/02/08 10:31:52 INFO mapreduce.Job: map 100% reduce 100%
15/02/08 10:31:53 INFO mapreduce.Job: Job job_1423408479621_0009 completed successfully
15/02/08 10:31:53 INFO mapreduce.Job: Counters: 49
```

<https://www.dropbox.com/s/nhz1cvcfteyoqlu/wordcount.jar>

:

```
$hadoop jar wordcount.jar org.myorg.WordCount
/user/cloudera/input/* /user/cloudera/output/wordcount
```

Reviewing MapReduce Job in Hue

HUE [Home](#) [Query Editors](#) [Data Browsers](#) [Workflows](#) [Search](#) [Security](#)

Job Browser

Username Text

Succeeded **Running** **Failed** **Killed**

Logs	ID	Name	Application Type	Status	User	Maps	Reduces	Queue	Priority	Duration	Submitted
	1465875170640_0001	wordcount	MAPREDUCE	SUCCEEDED	root	100%	100%	root.root	N/A	21s	06/13/16 21:32:37

Showing 1 to 1 of 1 entries

[← Previous](#) [1](#) [Next →](#)

Reviewing MapReduce Job in Hue

HUE [Home](#) [Query Editors](#) [Data Browsers](#) [Workflows](#) [Search](#) [Security](#)

Job Browser

JOB ID wordcount

146587517064...

TYPE

MR2

USER

root

STATUS

SUCCEEDED




LOGS

Logs

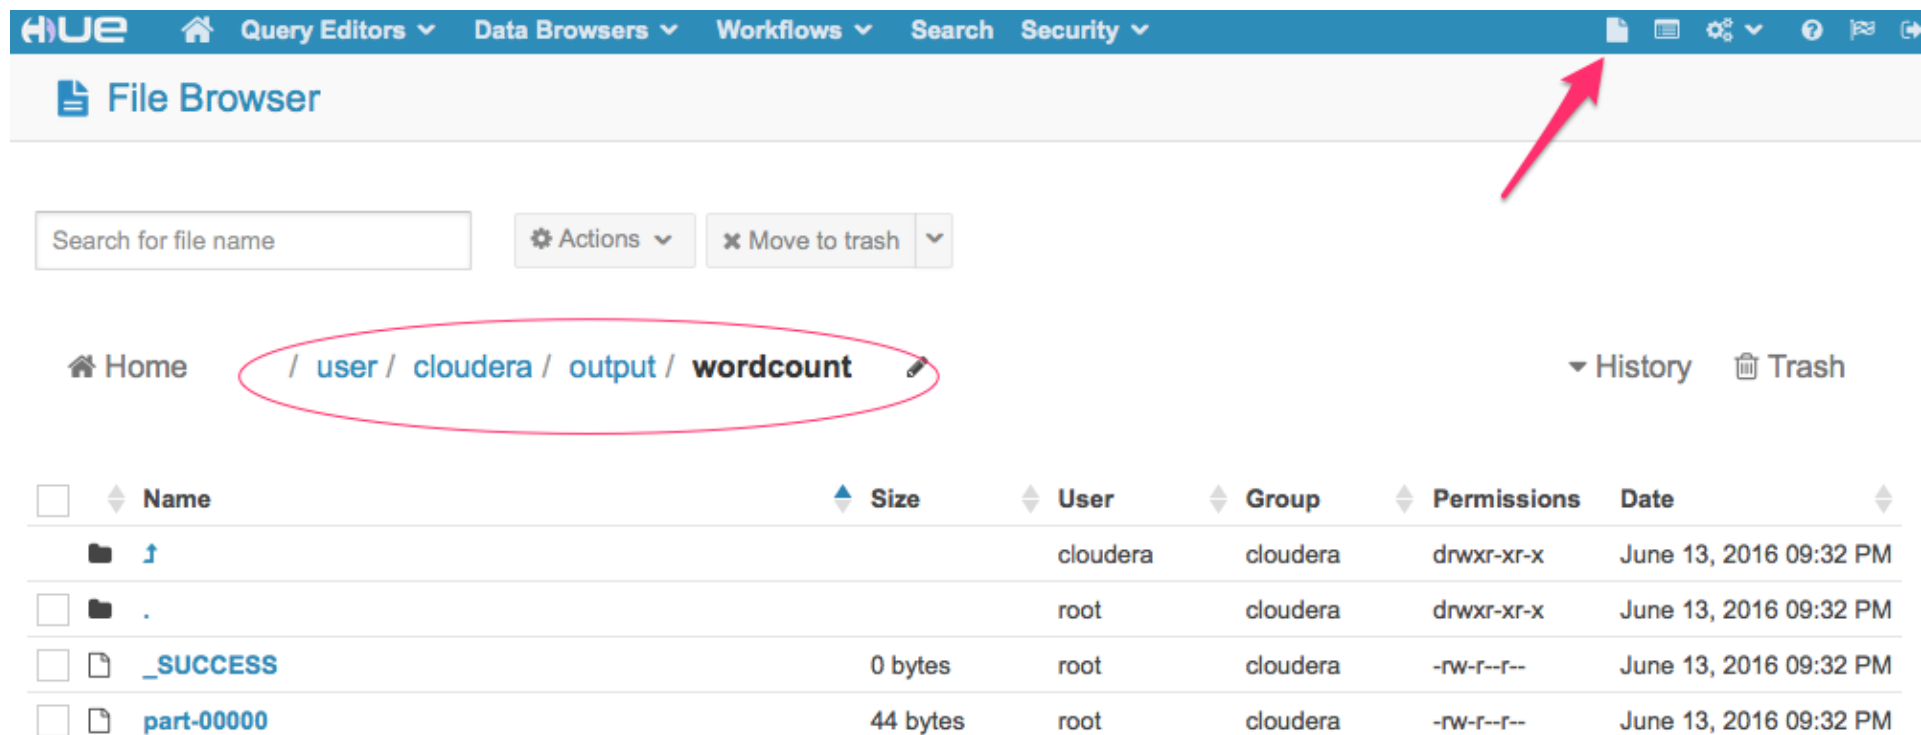
[Attempts](#) **Tasks** [Metadata](#) [Counters](#)

Recent Tasks

[View All Tasks »](#)

Logs	Tasks	Type
	task_1465875170640_0001_m_000000	MAP
	task_1465875170640_0001_m_000001	MAP
	task_1465875170640_0001_r_000000	REDUCE

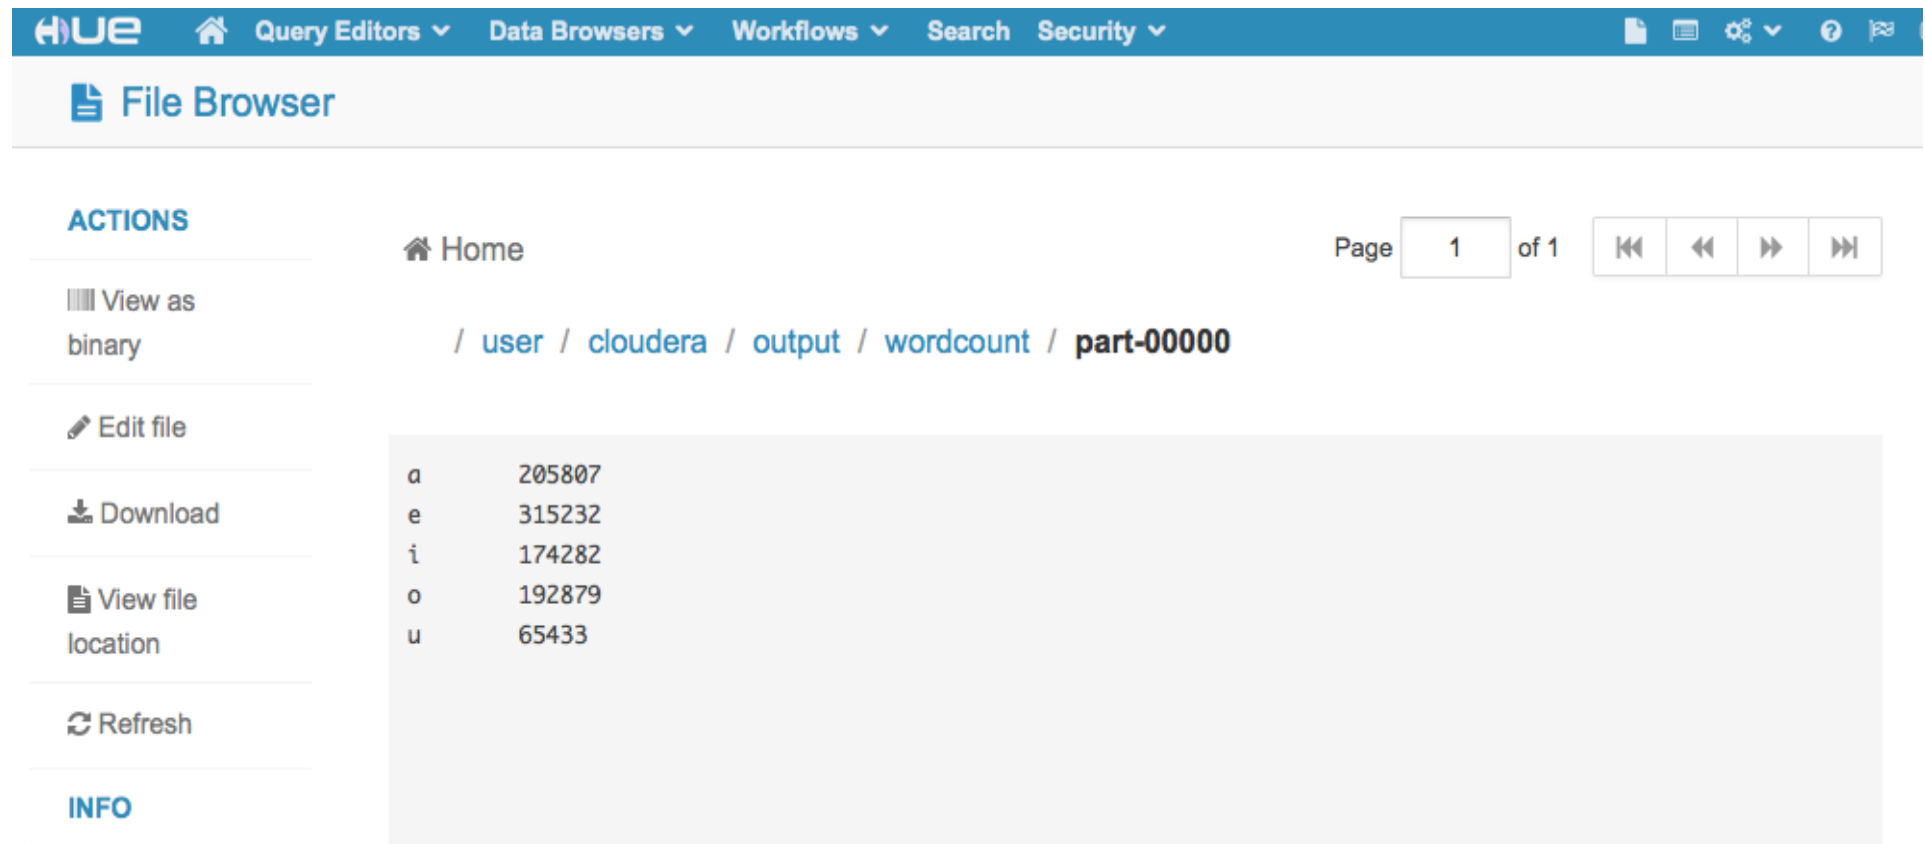
Reviewing MapReduce Output Result



The screenshot shows the Hue File Browser interface. The top navigation bar includes links for Query Editors, Data Browsers, Workflows, Search, and Security. A red arrow points to the top right corner of the interface. Below the navigation bar, the File Browser section shows a search bar and a list of actions. The current path is `/user/cloudera/output/wordcount`, which is highlighted with a red oval. The breadcrumb path is `Home / user / cloudera / output / wordcount`. The right side of the interface shows `History` and `Trash` links. Below the breadcrumb, there is a table listing files and directories.

Name	Size	User	Group	Permissions	Date
↑		cloudera	cloudera	drwxr-xr-x	June 13, 2016 09:32 PM
.		root	cloudera	drwxr-xr-x	June 13, 2016 09:32 PM
_SUCCESS	0 bytes	root	cloudera	-rw-r--r--	June 13, 2016 09:32 PM
part-00000	44 bytes	root	cloudera	-rw-r--r--	June 13, 2016 09:32 PM

Reviewing MapReduce Output Result



The screenshot displays the HUE File Browser interface. The top navigation bar includes the HUE logo and links for Query Editors, Data Browsers, Workflows, Search, and Security. Below this, the 'File Browser' section is active. On the left, a sidebar contains 'ACTIONS' (View as binary, Edit file, Download, View file location, Refresh) and 'INFO'. The main content area shows the breadcrumb path: / user / cloudera / output / wordcount / **part-00000**. The file content is displayed as a table with two columns: a character and its corresponding count.

a	205807
e	315232
i	174282
o	192879
u	65433



Lecture

Understanding Oozie

Introduction

Workflow scheduler for Hadoop



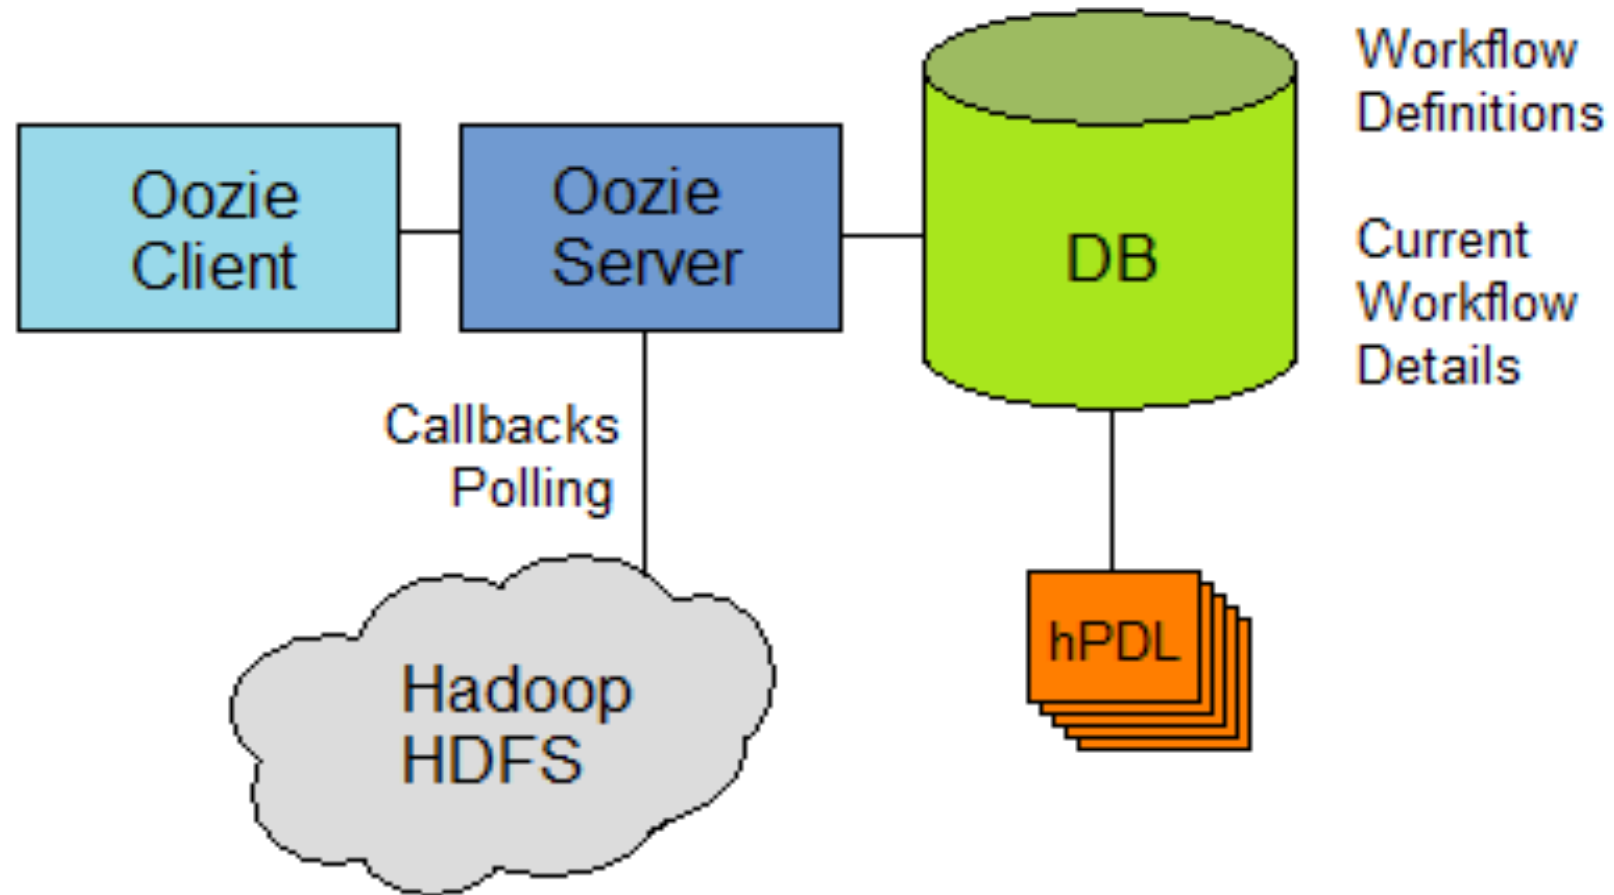
Oozie is a workflow scheduler system to manage Apache Hadoop jobs. Oozie is integrated with the rest of the Hadoop stack supporting several types of Hadoop jobs out of the box (such as Java map-reduce, Streaming map-reduce, Pig, Hive, Sqoop and Distcp) as well as system specific jobs (such as Java programs and shell scripts).



What is Oozie?

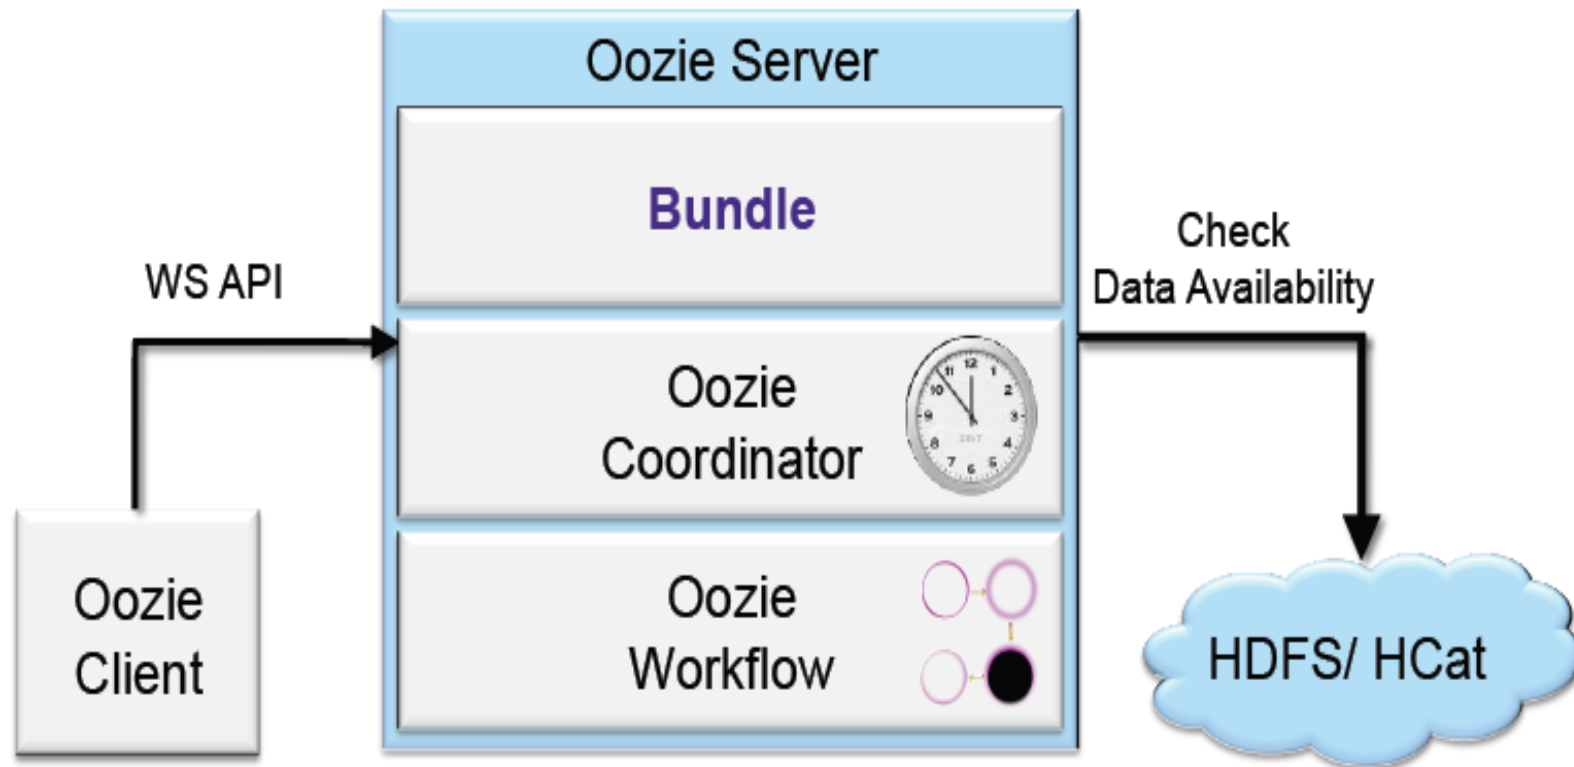
- Work flow scheduler for Hadoop
- Manages Hadoop Jobs
- Integrated with many Hadoop apps i.e. Pig, Hive
- Scalable
- Schedule jobs
- A work flow is a collection of actions.
- A work flow is
 - Arranged as a DAG (direct acyclic graph)
 - Graph stored as hPDL (XML process definition)

Oozie Architecture

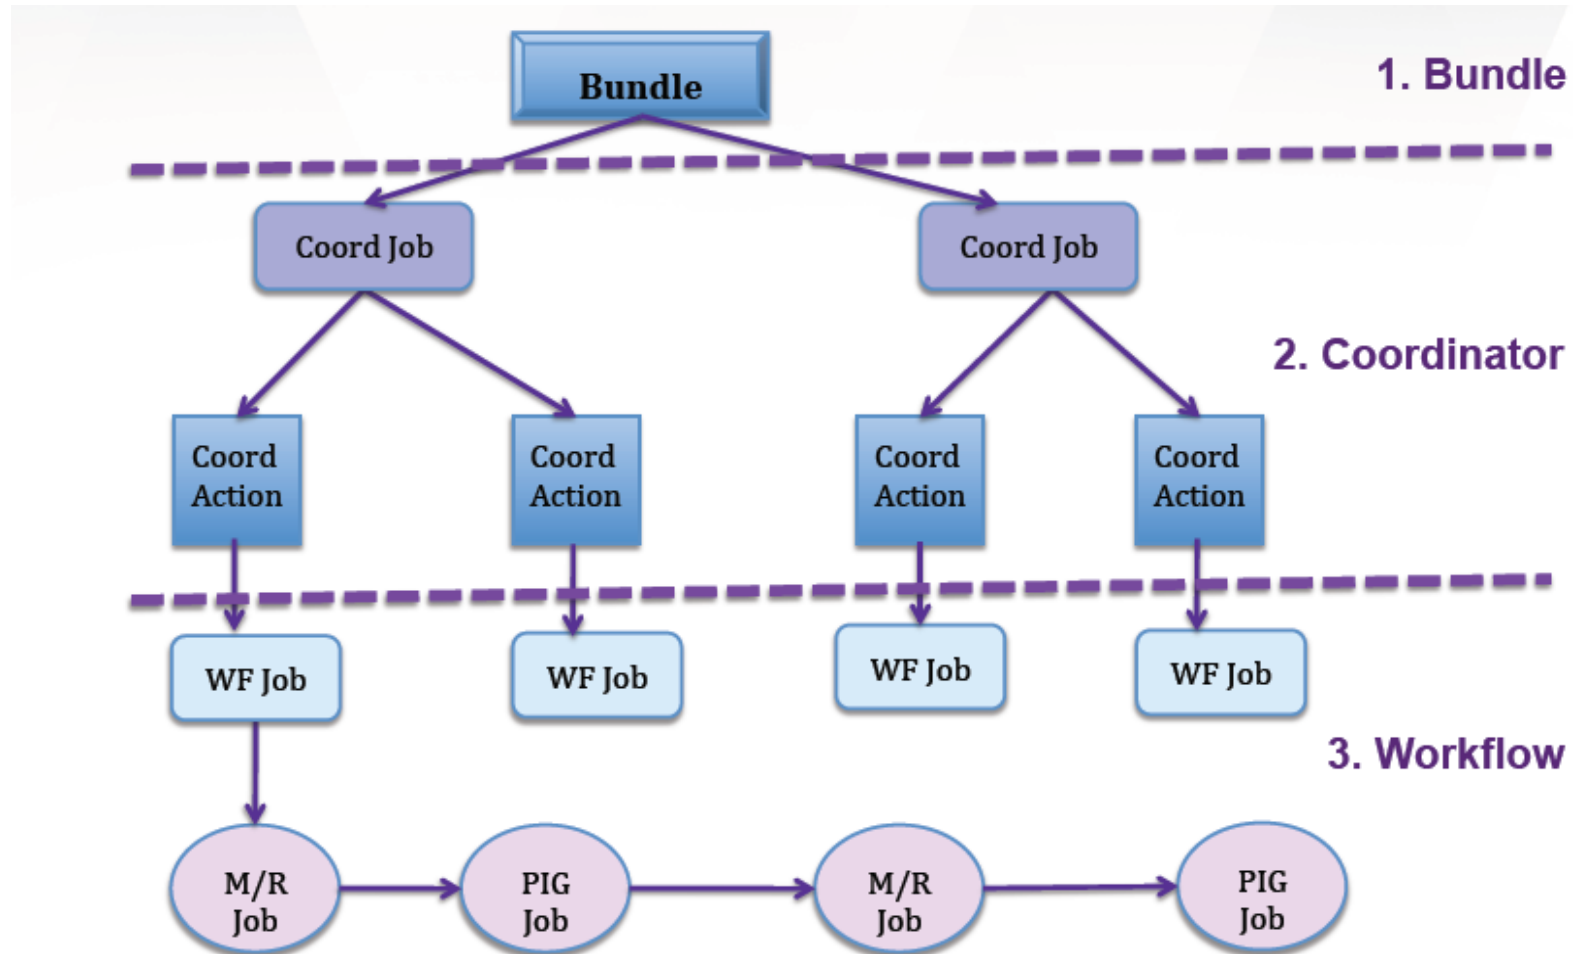


Source: info@semtech-solutions.co.nz

Oozie Server



Layer of Abstraction in Oozie

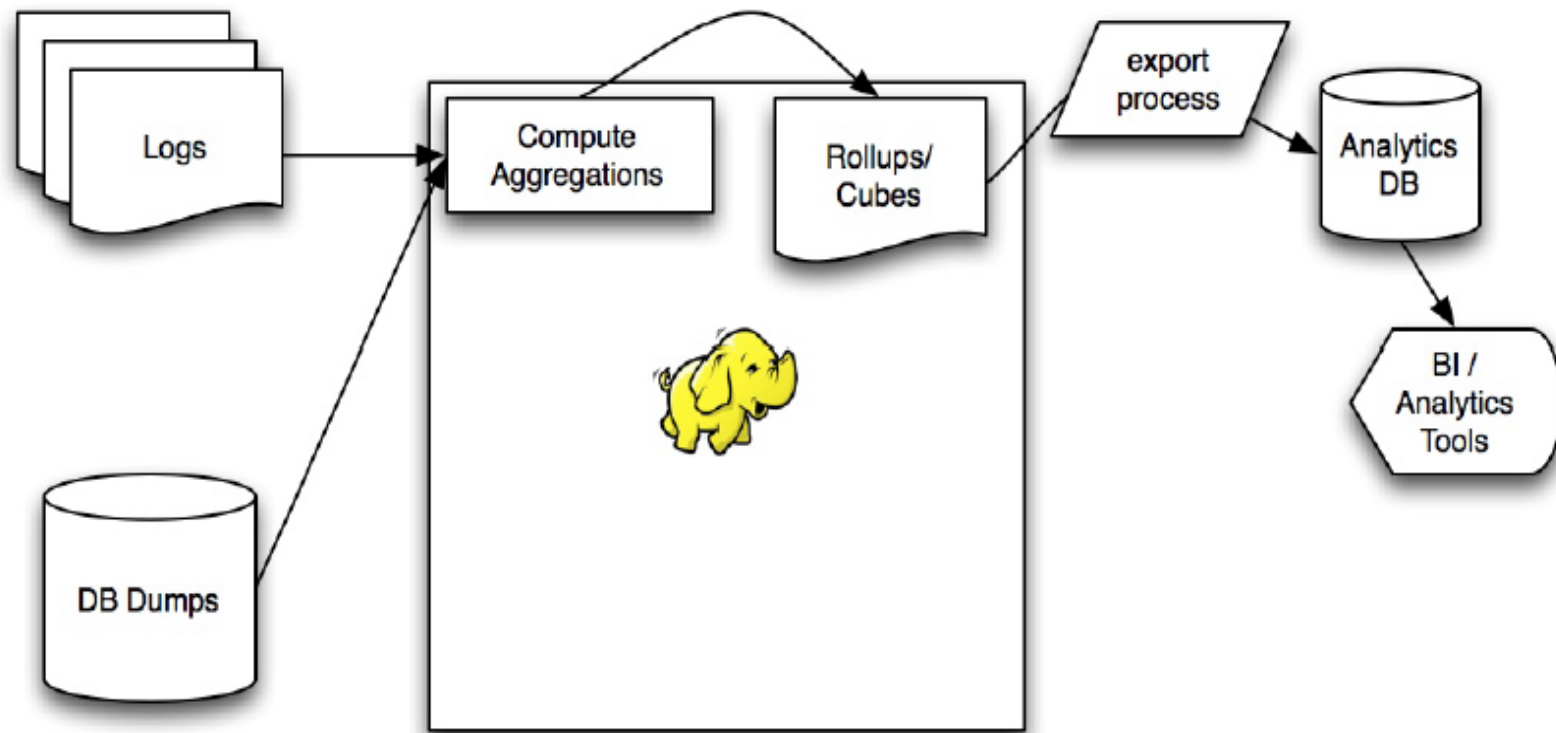




Workflow Example: Data Analytics

- Logs => fact table(s)
- Database backup => Dimension tables
- Complete rollups/cubes
- Load data into a low-latency storage (e.g. Hbase, HDFS)
- Dashboard & BI tools

Workflow Example: Data Analytics



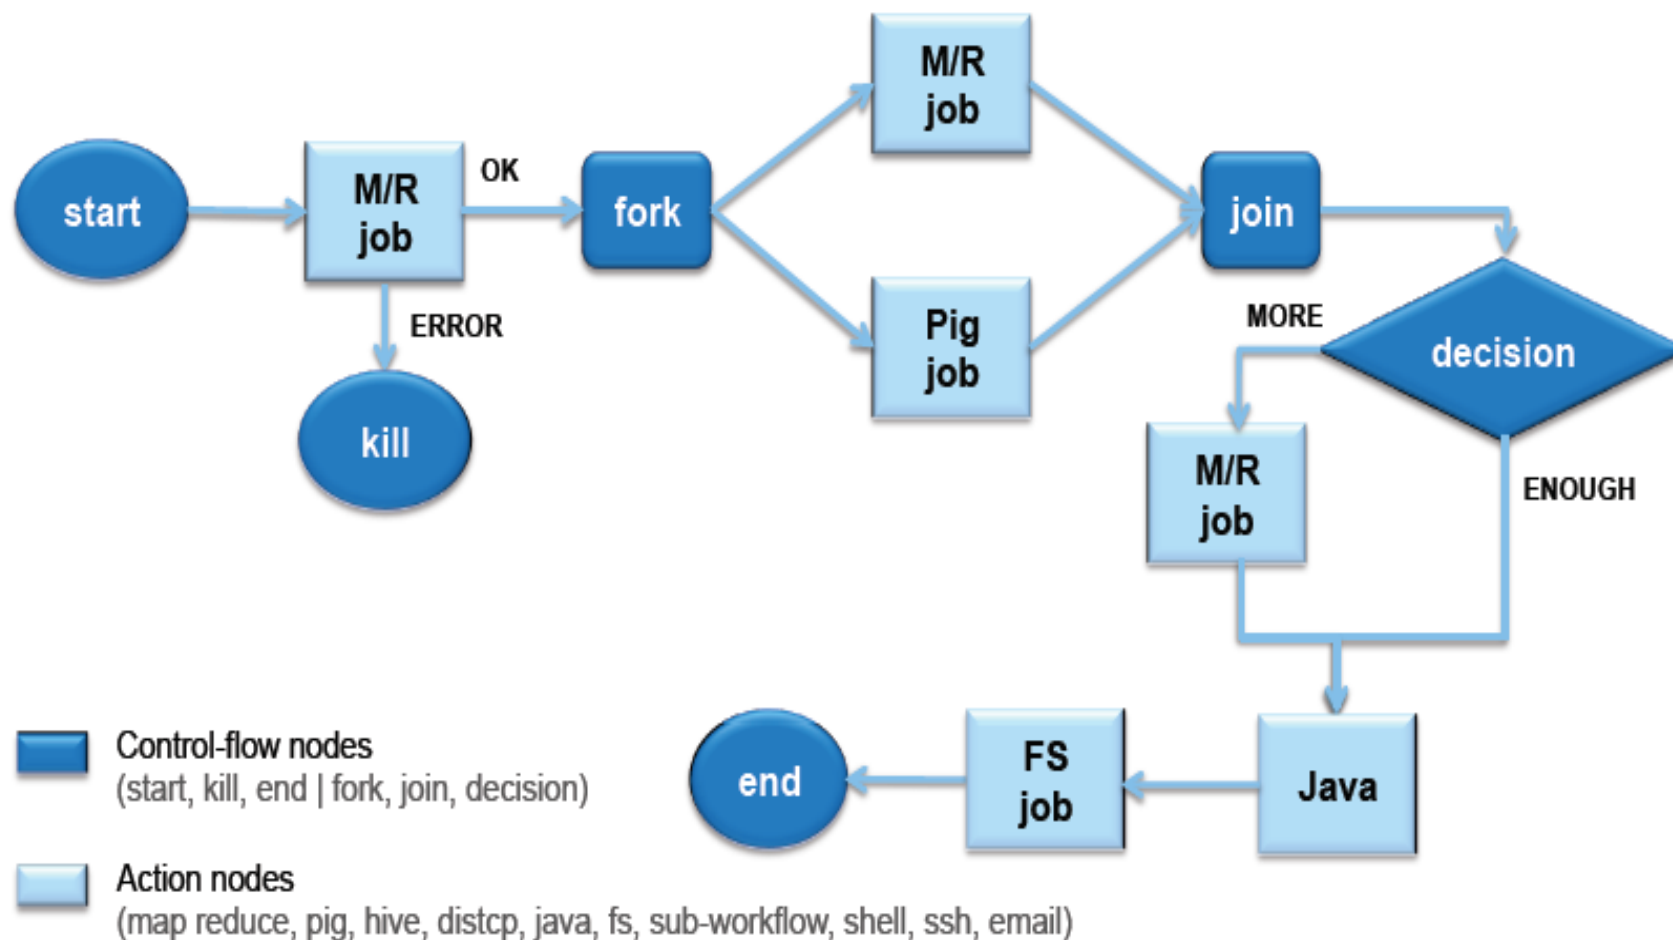
Source: Workflow Engines for Hadoop, Joe Crobak, 2013



Workflow Example: Data Analytics

- What happens if there is a failure?
 - Rebuild the failed day
 - .. and any downstream datasets
- With Hadoop Workflow
 - Possible OK to skip a day
 - Workflow tends to be self-contained, so you do not need to run downstream.
 - Sanity check your data before pushing to production.

Oozie Workflow



Source: Oozie – Now and Beyond, Yahoo, 2013



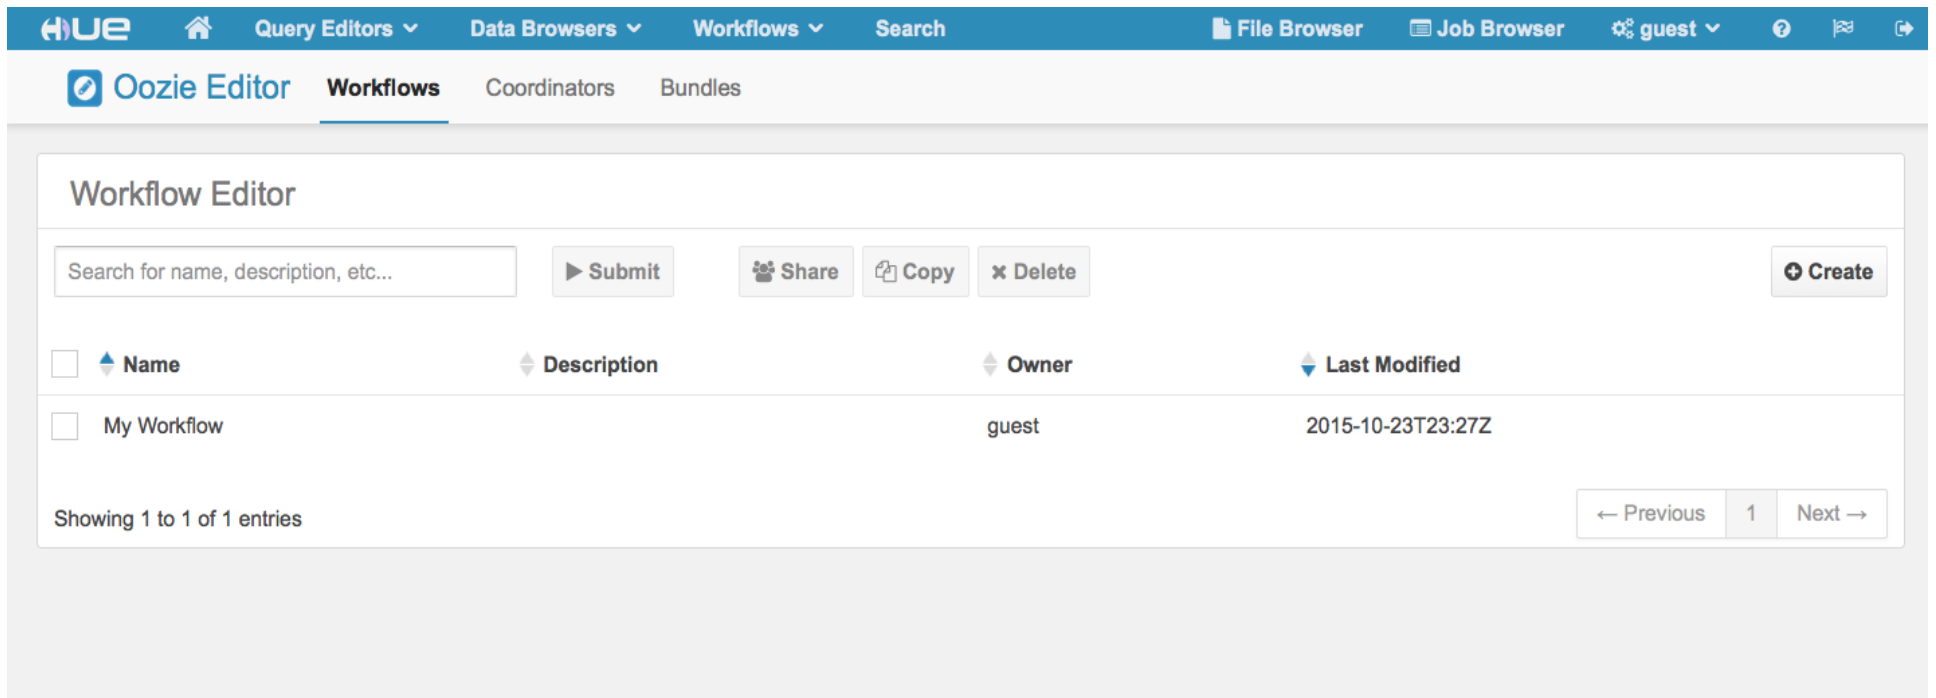
Oozie Use Cases

- Time Triggers
 - Execute your workflow every 15 minutes
- Time and Data Triggers
 - Materialize your workflow every hour, but only run them when the input data is ready (that is loaded to the grid every hour)
- Rolling Window
 - Access 15 minute datasets and roll them up into hourly datasets



Hands-On: Running Map Reduce using Oozie workflow

Using Hue: select WorkFlows >> Editors >> Workflows



The screenshot shows the Hue Oozie Editor interface. The top navigation bar includes links for Query Editors, Data Browsers, Workflows, Search, File Browser, Job Browser, and a user profile for 'guest'. Below this, the 'Oozie Editor' section has tabs for Workflows, Coordinators, and Bundles. The 'Workflows' tab is active, displaying a 'Workflow Editor' section. This section contains a search bar with the placeholder text 'Search for name, description, etc...', and buttons for Submit, Share, Copy, Delete, and Create. Below the search bar is a table with the following columns: Name, Description, Owner, and Last Modified. The table contains one entry: 'My Workflow', owned by 'guest', with a last modified date of '2015-10-23T23:27Z'. At the bottom left, it says 'Showing 1 to 1 of 1 entries', and at the bottom right, there are navigation buttons for 'Previous', '1', and 'Next'.

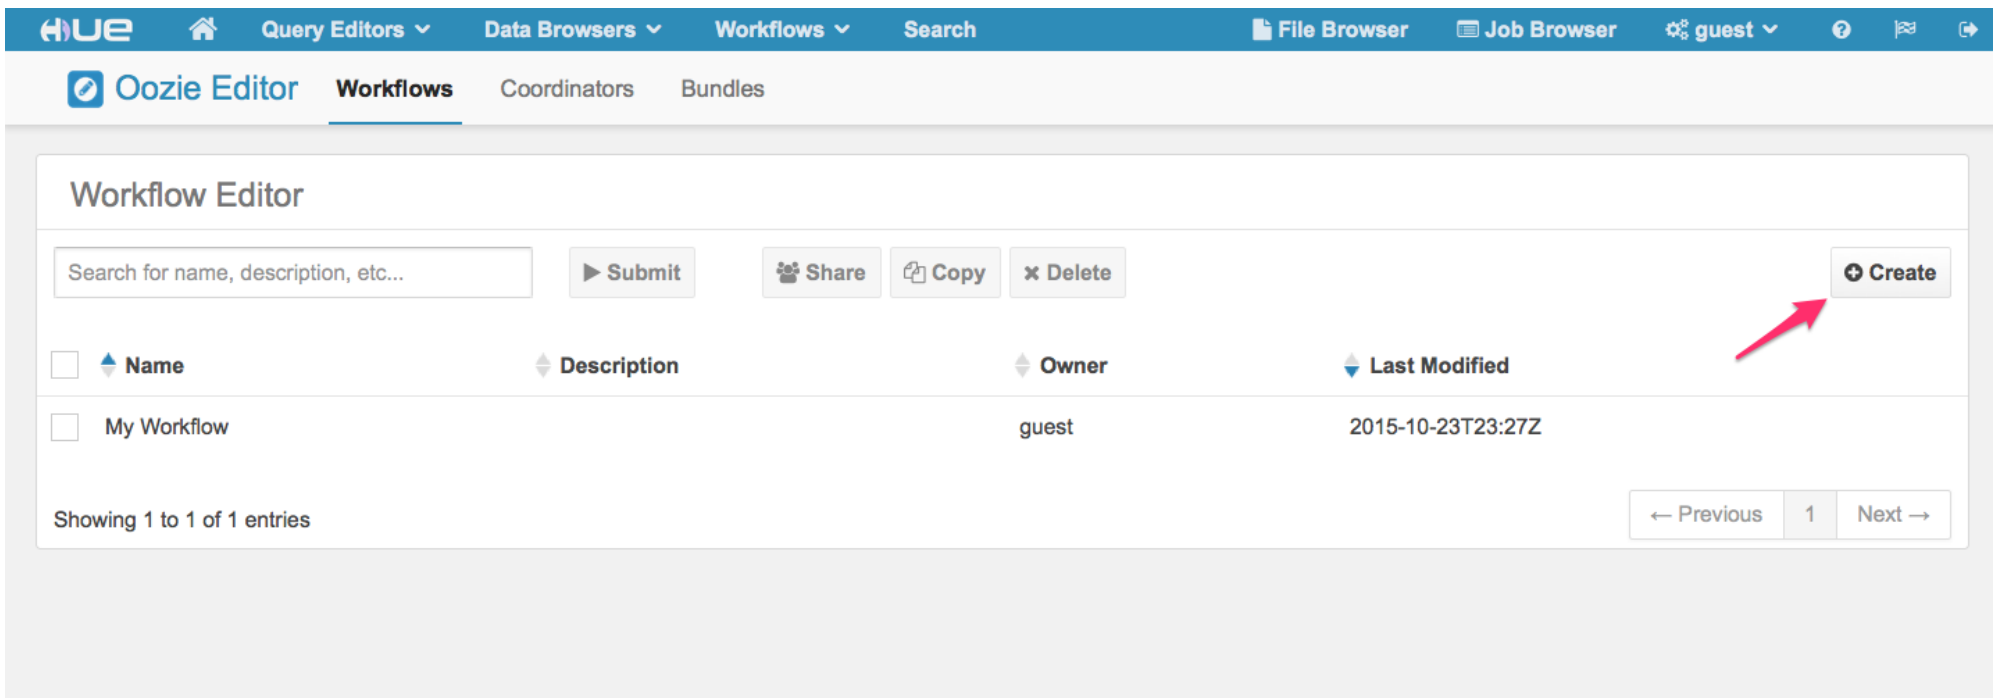
<input type="checkbox"/>	Name	Description	Owner	Last Modified
<input type="checkbox"/>	My Workflow		guest	2015-10-23T23:27Z

Showing 1 to 1 of 1 entries

← Previous 1 Next →

Create a new workflow

- Click Create button; the following screen will be displayed
- Name the workflow as WordCountWorkflow



The screenshot shows the Oozie Editor interface. The top navigation bar includes links for Home, Query Editors, Data Browsers, Workflows, Search, File Browser, Job Browser, and a user profile (guest). The main content area is titled 'Workflow Editor' and features a search bar and action buttons: Submit, Share, Copy, Delete, and Create. A table below lists existing workflows. A red arrow points to the 'Create' button.

<input type="checkbox"/>	Name	Description	Owner	Last Modified
<input type="checkbox"/>	My Workflow		guest	2015-10-23T23:27Z

Showing 1 to 1 of 1 entries

Navigation: ← Previous | 1 | Next →




HUE Home Query Editors ▾ Data Browsers ▾ Workflows ▾ Search File Browser Job Browser guest ▾ ?

Oozie Editor Workflows Coordinators Bundles Unsaved [edit icon] [settings icon] [trash icon] [save icon] [new icon]


ACTIONS [hand icon] [hand with 2 icon] [pig icon] [star icon] [code icon] [refresh icon] [file with drop icon] [link icon] [greater than icon] [phone icon] [document icon] [envelope icon] [double arrow icon] [stack of papers icon] [blue square icon]

WordCountWorkflow

Add a description...

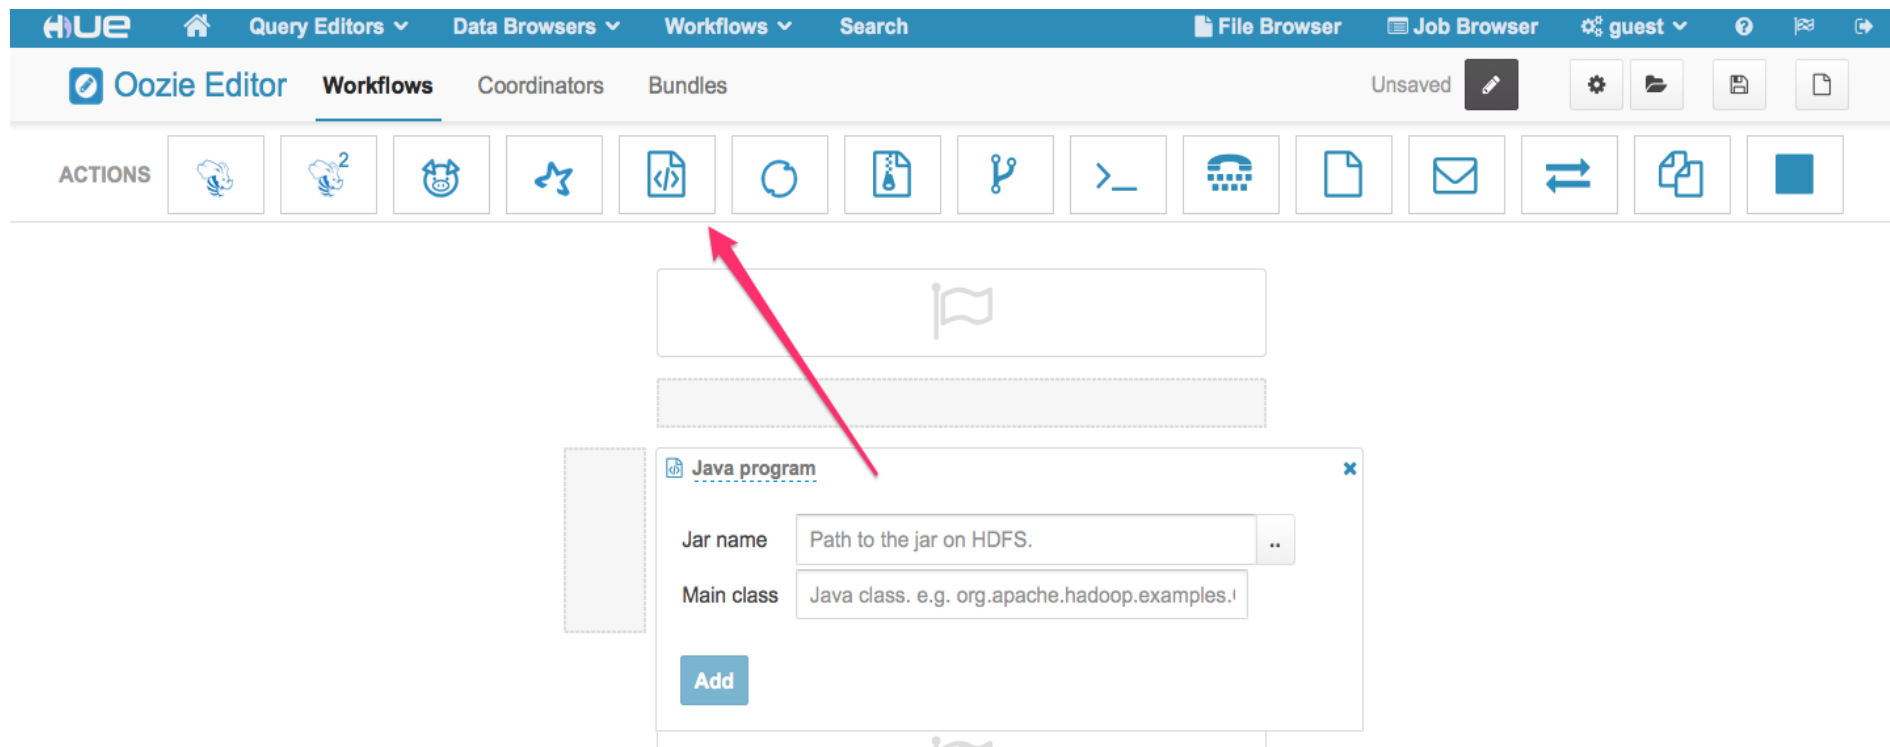


Drop your action here



Select a Java job for the workflow

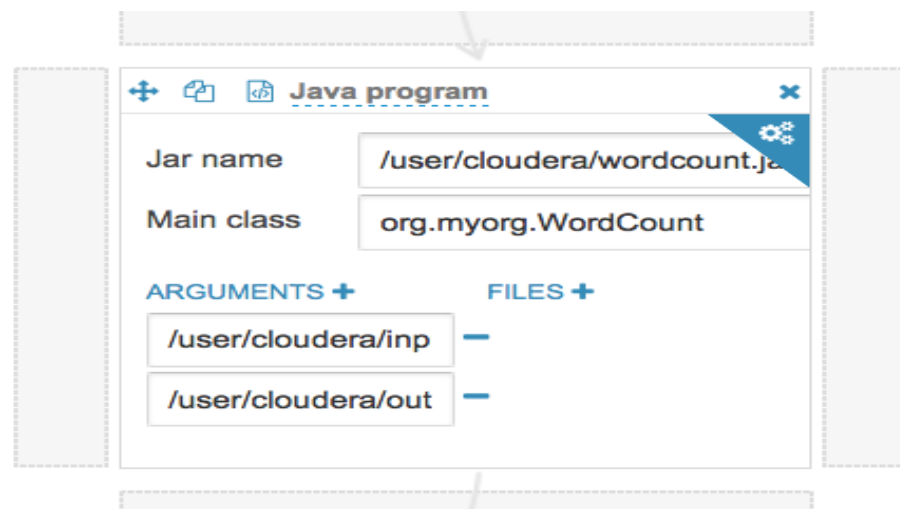
From the Oozie editor, drag **Java Program** and drop between start and end



Edit the Java Job

Assign the following value

- Jar name: wordcount.jar (select ... choose upload from local machine)
- Main Class: org.myorg.WordCount
- Arguments: /user/cloudera/input/*
- /user/cloudera/output/wordcount



The screenshot shows a configuration window titled "Java program" with a close button (X) in the top right corner. The window contains the following fields and sections:

- Jar name:** /user/cloudera/wordcount.jar
- Main class:** org.myorg.WordCount
- ARGUMENTS +:** A section with two input fields: /user/cloudera/inp and /user/cloudera/out.
- FILES +:** A section with two input fields: /user/cloudera/inp and /user/cloudera/out.

Submit the workflow

- Click Done, follow by Save
- Then click submit

The screenshot displays the HUE Oozie Editor interface. The top navigation bar includes links for Query Editors, Data Browsers, Workflows, Search, File Browser, Job Browser, and a user dropdown menu. The main toolbar contains icons for editing, running, saving, and other actions. Two red arrows point to the 'Submit' (play button) and 'Save' (floppy disk icon) buttons in the toolbar, with the labels 'Submit' and 'Save' written in red below them. Below the toolbar, a workflow diagram is visible, showing a sequence of steps. The first step is a 'Java program' action, which is expanded to show its configuration. The configuration includes a 'Jar name' field with the value '/user/guest/wordcount.jar', a 'Main class' field with the value 'org.myorg.WordCount', and two 'ARGUMENTS' fields: 'input/*' and 'output/wordcount'.